

Theodoros Vasiloudis

CONTACT INFORMATION

SICS Swedish ICT AB
Decisions, Networks and Analytics Lab
Swedish Institute of Computer Science
Royal Institute of Technology, KTH

Phone: +46(0)700477423
tvas@sics.se
tvas@kth.se

RESEARCH SUMMARY

My research deals with large scale learning from graphs and machine learning, including the development of novel algorithms and contributions to data analytics frameworks.

The projects that I am involved in include developing algorithms that make use of frameworks such as Apache Spark and Apache Flink, and contributing to the design and implementation of learning systems capable of handling massive datasets, such as the machine learning library for Apache Flink. As a committer for Apache Flink, I am responsible for coordinating the machine learning efforts for the project.

A research area that I am currently exploring is online learning, and more specifically distributed online boosting for decision trees. This year I will also be helping in the development of a distributed online/streaming machine learning library.

EDUCATION

Royal Institute of Technology, Stockholm, Sweden

Ph.D. Candidate

November 2014 – Currently

- Academic Supervisor: Professor Anders Holst
- Academic Co-supervisor: Professor Seif Haridi
- Industrial Supervisor: Daniel Gillblad

Royal Institute of Technology, Stockholm, Sweden

M.Sc. in Machine Learning

August 2012 – July 2014

- Thesis: Extending recommendation algorithms by modelling user context.
- Performed at: Spotify AB
- Supervisor: Associate Professor Hedvig Kjellström
- Industrial Supervisor: Boxun Zhang

Aristotle University of Thessaloniki, Thessaloniki, Greece

B.Sc. in Computer Science

September 2004 – March 2011

- Thesis: Implementation of two algorithms for the classification of multi-label data using Bayesian networks.
- Supervisor: Assistant Professor Grigorios Tsoumakas

PUBLICATIONS

Theodoros Vasiloudis, Puya Vahabi, Ross Kravitz, Valery Rashkov. *Predicting Session Length in Media Streaming*.

ACM International Conference on Research and Development in Information Retrieval (SIGIR) 2017, Tokyo, Japan.

Olof Görnerup, Daniel Gillblad, **Theodoros Vasiloudis**. *Domain-Agnostic Discovery of Similarities and Concepts at Scale*.

Knowledge and Information Systems, Springer (2016).

Olof Görnerup, Daniel Gillblad, **Theodoros Vasiloudis**. *Knowing an Object by the Company It Keeps: A Domain-Agnostic Scheme for Similarity Discovery*.

IEEE International Conference on Data Mining (ICDM) 2015, Atlantic City, USA.

PROFESSIONAL EXPERIENCE

Swedish Institute of Computer Science, Stockholm, Sweden

Researcher

August 2014 – Currently

I am currently employed as a researcher at the Swedish Institute of Computer Science (SICS) where I also have the opportunity to pursue a PhD title from the Royal Institute of Technology, KTH. My work as a researcher is aligned to my research interests in large scale machine learning algorithms and applications.

Pandora Media, Oakland, CA, USA

Scientist intern

June 2016 – August 2016

During the summer of 2016 I was a scientist intern at Pandora, working on user modelling and more specifically predicting the length of user sessions on the service. Utilizing techniques from survival analysis, mixed effects modeling and gradient boosted trees we were able to significantly outperform previous approaches and produce new insights on user's use of the service.

Data Artisans, Berlin, Germany

Software Engineering Internship

April 2015 – July 2015

During my time at Data Artisans I contributed to the design and implementation of the Machine Learning library for Apache Flink. This included work on algorithms and tooling, such as developing a convex optimization framework, aiding in the design of machine learning pipelines, and developing an evaluation and cross-validation framework for machine learning models. I also assisted in expanding the library by reviewing and advising committers and students from Google Summer of Code and TU Berlin.

Spotify, Stockholm, Sweden

Master Thesis

January 2014 – July 2014

My thesis was performed at Spotify AB, where I studied the effects of including user context in the recommendation pipeline. We explored the effects of adding contextual information like the day of the week into a traditional recommender system, and showed improvements in performance vs. baselines like item-to-item similarity.

GITHUB PROFILE <http://github.com/thvasilo>

PROJECTS
COMPLETED

American Epilepsy Society Seizure Prediction Challenge

Theodoros Vasiloudis, Johan Tjellén, Erik Ylipää

The task for this competition was the prediction of oncoming seizures using multivariate EEG signals as input. Using features such as channel cross-correlation and employing SVM classification our team achieved a top 30 finish, out of the 500 teams participating.

Improving Search Replies of the Spotify Search Engine

Theodoros Vasiloudis, Julius Adorf, Caroline Kéramsi, Mikaela Nöteberg

In this project we were tasked with improving the performance of the spelling correction system at Spotify. A Bayesian approach was employed, using an error model and a language model learned from observed user data.

SKILLS

Programming: Scala, Python, Java, C++

Scientific computing: R, SciPy stack, MATLAB, CUDA

Big data: Apache Spark, Apache Flink, Hadoop MapReduce, Turi Create

Data visualization: ggplot2, Matplotlib