

Approximate queries and graph streams on Flink

Theodore Vasiloudis, Swedish Institute of Technology

tvas@kth.se
[@thvasilo](#)

Hello, I'm Theo!



**RI.
SE**



dataArtisans

pandora

amazon

Motivation

- We want analytics that provide us with immediate answers
 - *Hadoop: Yesterday's insights, tomorrow!*
- Data never stops!
 - Infinite memory?
- Solution: windows and approximations
 - Windows give us a snapshot of the world
 - Approximations allow us to continuously measure the world
 - First part of presentation is about approximate streaming algorithms, second focuses on using windows for graph analytics

Approximate Queries on Flink

Work by Tobias Lindener, KTH

<https://github.com/tlindener/ApproximateQueries/>

End Goal: Approximate SQL queries on Flink

```
SELECT avg(sessionTime)  
FROM Table  
WHERE city='San Francisco'  
WITHIN 2 SECONDS
```

Queries with Time Bounds

```
SELECT avg(sessionTime)  
FROM Table  
WHERE city='San Francisco'  
ERROR 0.1 CONFIDENCE 95.0%
```

Queries with Error Bounds

First step: Sketches for standing queries

Web Site Logs

Time	User ID	Site	Time Spent Sec	Items Viewed
9:00 AM	U1	Apps	59	5
9:30 AM	U2	Apps	179	15
10:00 AM	U3	Music	29	3
1:00 PM	U1	Music	89	10
Billions of rows ...				

Financial Transactions System Log

Time	User ID	Site	Purchased	Revenue
9:00 AM	U1	Apps	FaceTune	\$3.99
9:30 AM	U2	Apps	Minecraft	\$6.99
10:00 AM	U3	Music	Purple Rain	\$1.29
Billions of rows ...				

- Num. unique users who visited both Apps and Music over the last hour
- Median and 95%ile Time Spent over the last day
- Most frequently purchased songs

Sketch Algorithms for Massive Data


- Research area since the 70s (Knuth, Indyk, Flajolet)
- Goal: Efficient (compute+memory) algorithms for “simple” tasks
 - Frequent items
 - Set cardinality
 - Moments (mean, median, variance etc.)
 - Quantiles and histograms
 - Graph algorithms (triangle count, connected components)
 - Nearest neighbors
- We use them as building blocks for more complex algorithms
 - Databases (joins)
 - Machine learning (decision trees)

Yahoo DataSketches


- Highly optimized sketch library
- Apache Licensed
- Available for Pig, Hive



Big Win #1: Size of the Query Process



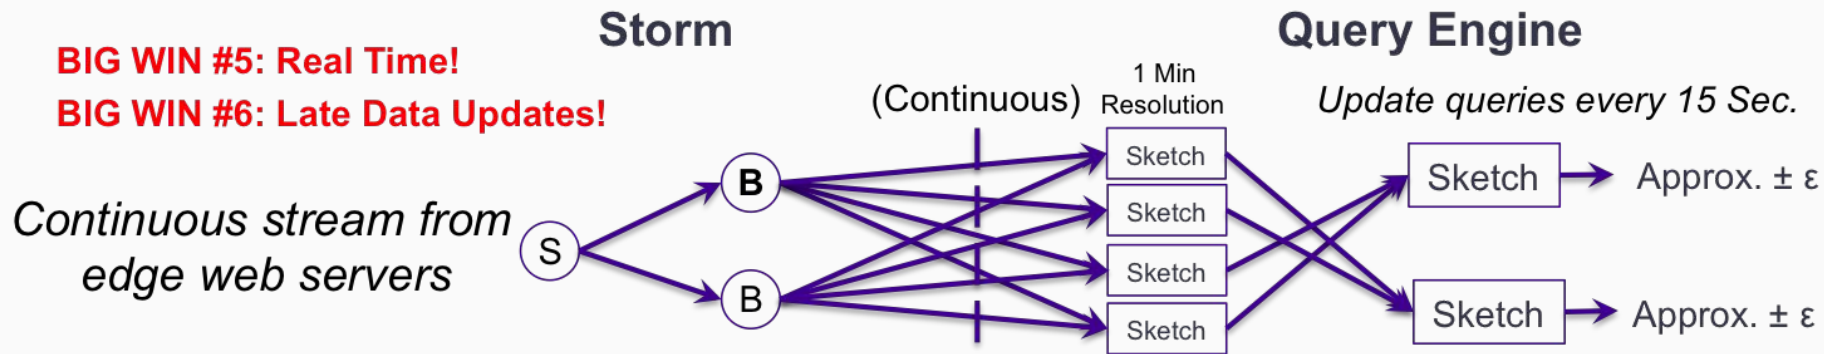
Big Win #2: Sketch Mergeability Enables Parallel Processing



Big Wins #3 & 4: Query Speed, Architecture Simplicity



Big Wins #5 & 6: Real Time, Late Data Updates

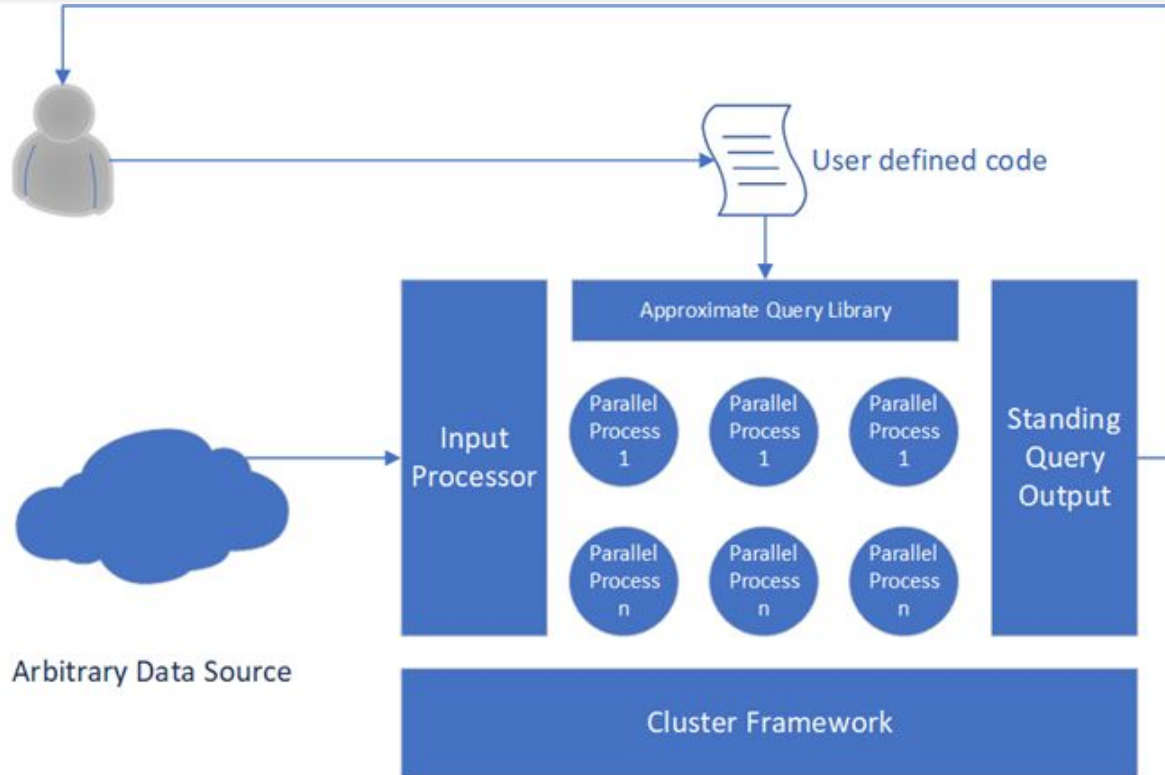


Why Flink?

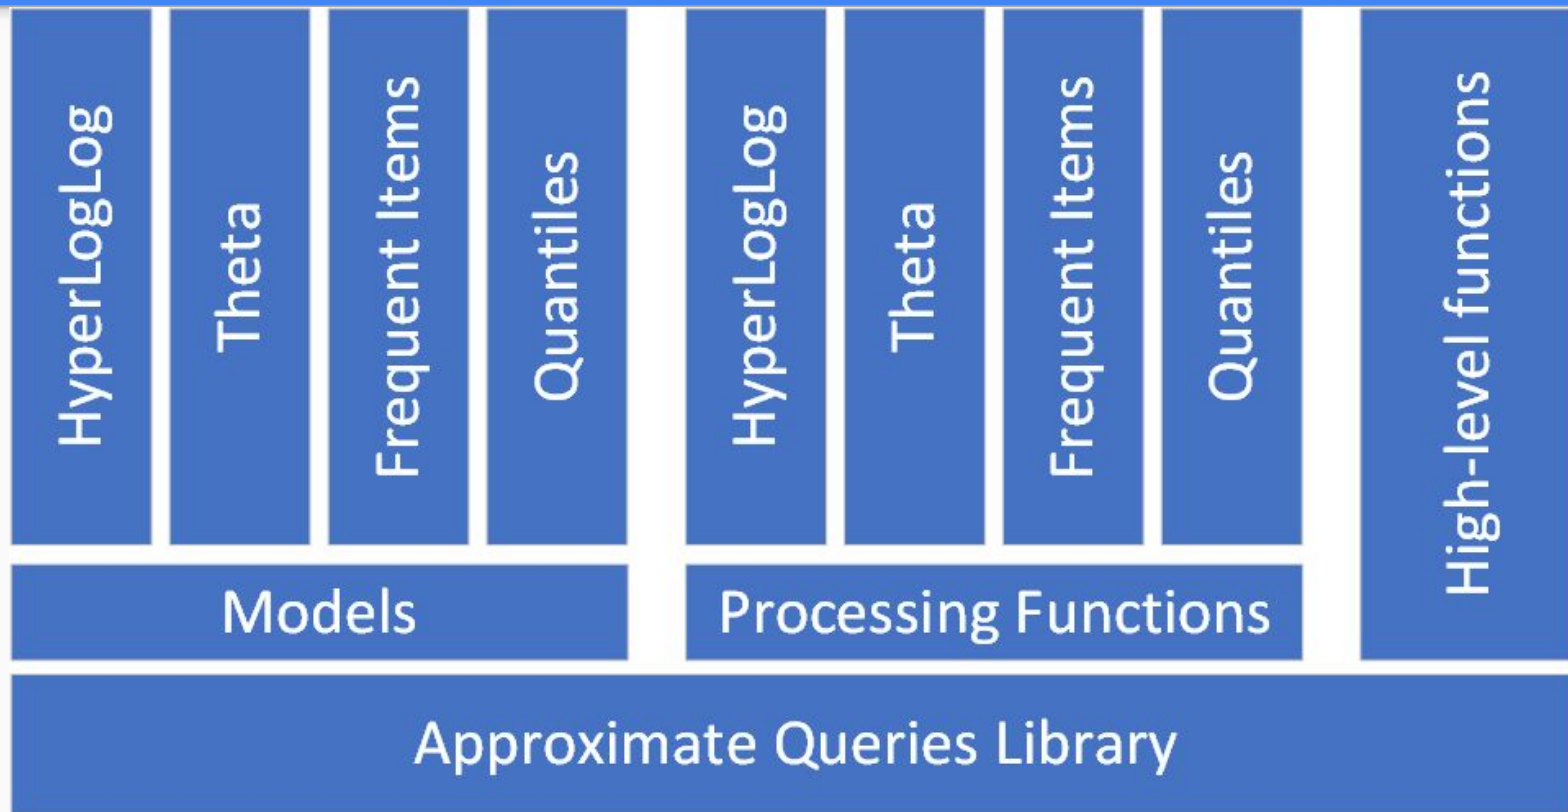
- From PR sketches-core#81: *Sketch now implements Serializable* (closed)
- Lee Rhodes (package author) writes:

*Sketches are streaming algorithms and are **stateful by design**. Attempting to force them into a stateless paradigm will result in orders-of-magnitude poorer performance. It is like pounding a square peg into a round hole.*

Library Design



Library Design



Query API

- Cardinality Estimation Queries

```
1 public static <T> DataStream<HllSketchAggregation>  
  ↳ runContinuousHll(DataStream<T> input, KeySelector  
  ↳ keySelector, KeySelector valueSelector, int emitMin)  
2  
3 public static <T> DataStream<ThetaSketchAggregation>  
  ↳ runContinuousTheta(DataStream<T> input, KeySelector  
  ↳ keySelector, KeySelector valueSelector, int emitMin)
```


Query API


- Frequent Items & Quantiles

```
1 public static <T> DataStream<TopNQueryResult>  
    ↳ continuousFrequentItems(DataStream<T> inputStream,  
    ↳ KeySelector valueSelector, int maxItems, int  
    ↳ emitMin)  
2  
3 public static <T> DataStream<QuantileQueryResult>  
    ↳ continuousQuantiles(DataStream<T> inputStream,  
    ↳ KeySelector valueSelector)
```

Experiments


Datasets


- Amazon Reviews
 - ~84M reviews of Amazon products
 - User, Item, Rating, Timestamp
- WikiTrace
 - Wikipedia access logs
 - ~80M requests, ~7M URLs



Runtime and memory consumption

	Runtime	Memory
Exact	~700s	~11GB
Sketch	~90s	~2GB





Amazon: Accuracy of frequent item sketch for different map sizes (lower is better)

Summary

- Built library to allow to easily run approximate queries on top of Flink
- End goal is to enable approximate streaming queries with SQL syntax

Gelly Stream: Streaming Graph Processing

Paris Carbone, KTH, Flink Committer
Vasiliki Kalavri, ETH, Flink PMC


Motivation

- Graphs are powerful representations of many interactions
 - Social networks
 - Purchases
 - Media views
- Again, data are massive, constantly arriving, and unbounded
- So we need distributed streaming graph processing

Previous work

- Graph snapshots
 - Pregel, Giraph, GraphX
- Graph streams
 - Summaries, approximate algorithms
 - Semi-streaming (disk)

Load-Compute-Store




1. **Load** snapshot to memory
2. **Compute** state/superstep
3. **Store** updated graph state
4. Goto 1

Load-Compute-Store

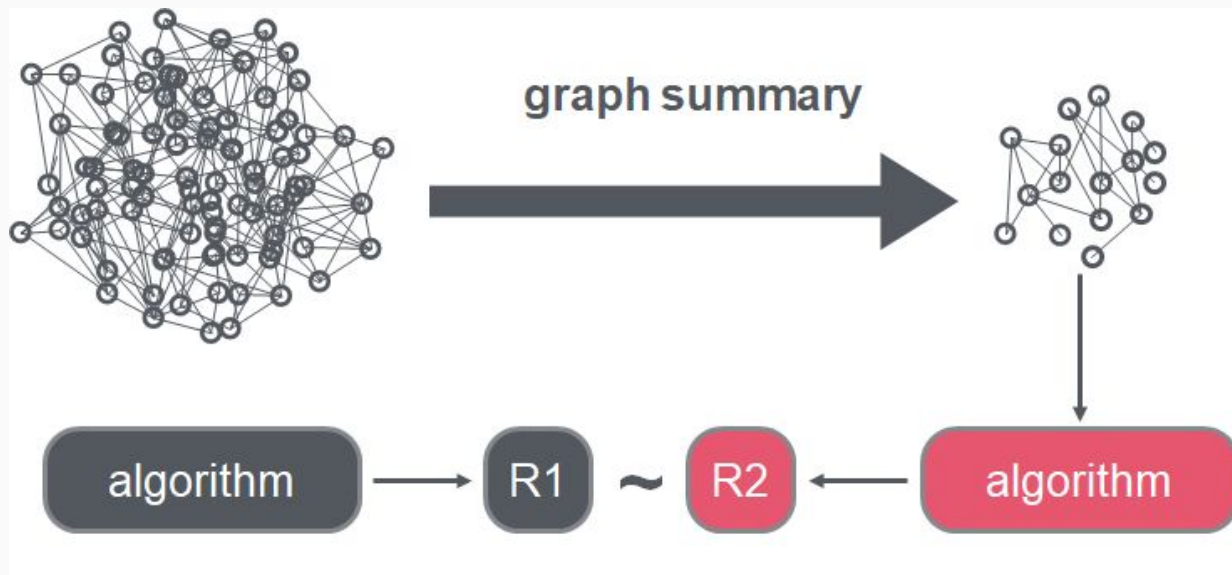
- Wide adoption: Pregel, Graphlab, GraphX
- Interfaces well with existing batch systems
- But has model issues:
 - Unnecessary **latency** for all graph measures.
 - Inefficient for incorporating **updates**
 - Sensitive to the **partitioning** method
 - **Re-computation** across snapshots

Graph Summaries: Intuition



Graph Summary Flavours

- **Spanners** : distance estimation
- **Sparsifiers** : cut estimation
- **Sketches** : homomorphic properties



Engineering benefits of stream processing

- **Low latency** and **high-throughput**
- **Long-running** processes can now pipeline computation
- Production Ready: end-to-end **fault tolerance**

Realizations brought by stream processing

1. Duality of **input** data + computational **state**
2. **Out-of-order** processing

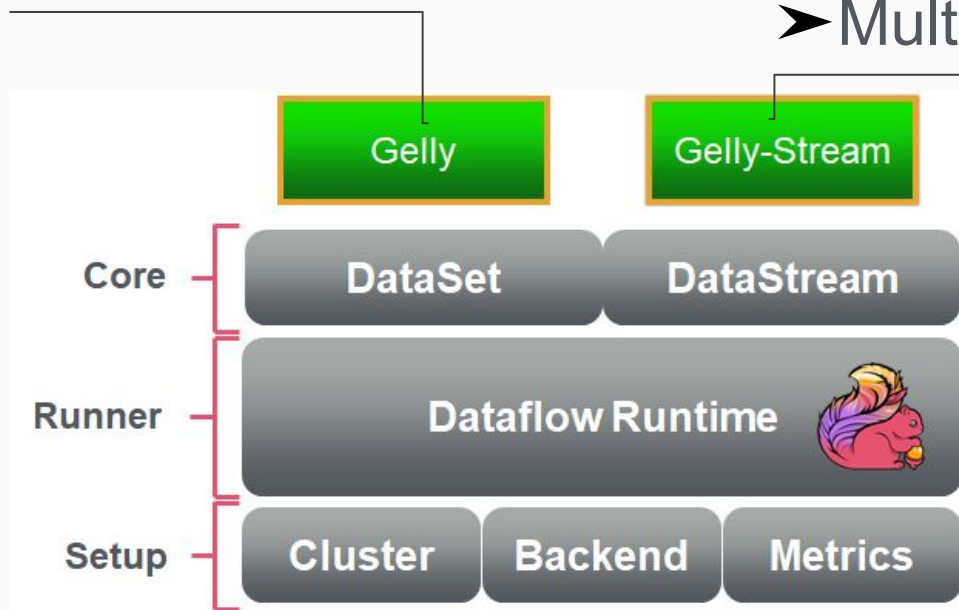
Exploiting stream processing for graphs

1. Duality of **input** data + computational **state**
 - a. Define evolving graph properties
 - b. Graph updates (input) \Leftrightarrow properties (state)
2. **Out-of-order** processing
 - a. Pre-compute blocking graph operations
 - b. Multiplex processing per snapshot or window

Gelly-stream overview

- Static Graphs
- Multi-Pass Algorithms
- Single Answer

- Dynamic Graphs
- Single-Pass
Properties/Summaries
- Multi-Pass on Snapshots



Gelly-Stream Data Types

- **EdgeStream** -> **Non-Blocking / Single-Pass Computation**
 - A distributed *data stream* consisting of **graph edge additions**.
 - Edges can contain **state** (e.g. weights).
 - Supports **property** streams, **transformations** and **aggregations**.
- **SnapshotStream** -> **Blocking / Multi-Pass Computation**
 - Each Snapshot is bounded~ i.e., static graph window.
 - It enables **neighborhood aggregations, iterations (e.g., BSP)**

EdgeStream Operations

Property Streams

EdgeStream \rightarrow DataStream

- `.getEdges()`
- `.getVertices()`
- `.numberOfVertices()`
- `.numberOfEdges()`
- `.getDegrees()`
- `.inDegrees()`
- `.outDegrees()`

Transformation Streams

EdgeStream \rightarrow EdgeStream


- `.mapEdges();`
- `.distinct();`
- `.filterVertices();`
- `.filterEdges();`
- `.reverse();`
- `.undirected();`
- `.union();`

EdgeStream Summaries

```
edgeStream.aggregate  
(new Summary(window, fold, combine, lower))
```


edgeStream.aggregate

(new Summary(window, fold, combine, lower))




edgeStream.aggregate

(new Summary(window, fold, combine, lower))





```
edgeStream.aggregate
```

```
(new Summary(window, fold, combine, lower))
```



edgeStream.**aggregate**

(new **Summary**(window, fold, combine, lower))



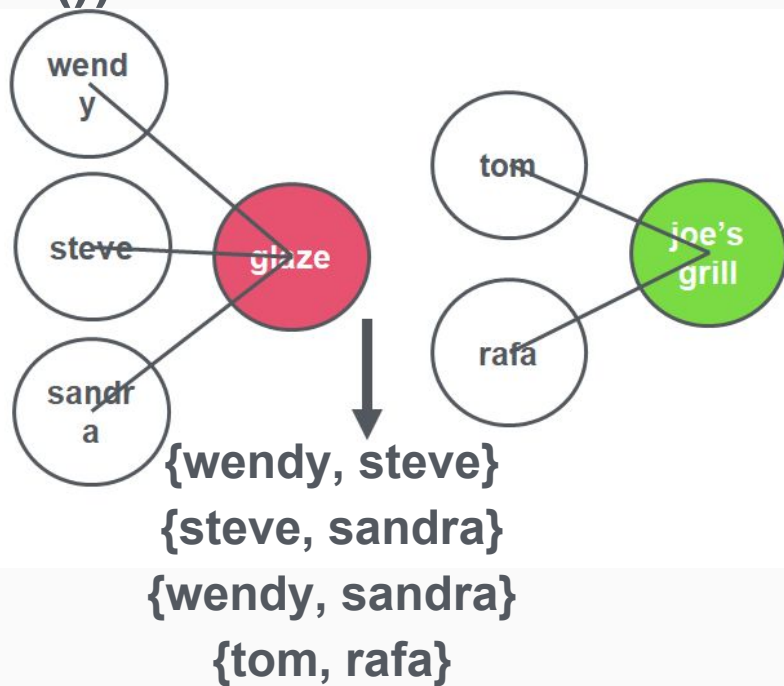
Provided Aggregates/Summaries

- Connected Components
- Bipartiteness Check (Binary)
- Window Triangle Count
- Rolling Triangle Count (Approximate)
- Continuous Degree Aggregate

Neighborhood Aggregation Example

```
edgeStream.filterVertices(DataScientists())  
.slice(Time.of(10, MINUTE), EdgeDirection.IN)  
applyOnNeighbors(FindPairs())
```

wendy checked_in **glaze**
steve checked_in **glaze**
tom checked_in **joe's_grill**
sandra checked_in **glaze**
rafa checked_in **joe's_grill**




Snapshot Iterations

- Most “deep” graph properties require **multiple passes**
- Sensitivity to **synchrony** during iterative processing depends on the algorithm and should be flexible (e.g., as in GraphLab).
- Avoiding scheduling delays (e.g. scheduling DataSet Iterations) is crucial for continuous processing.

Flink Stream Iterations

- A logical+physical loop redesign on Flink
- Introduces scoping and custom progress tracking
- Extends out-of-order dataflow processing
- Fully decentralised iterative execution

- ✦ termination condition
- ✦ level of (a)synchrony
- ✦ triggers



Take home message

- Streaming means unbounded
- Input \Leftrightarrow State
- Sketches and summaries let you deal with the unbounded nature of data using limited resources

Thank you!

tvas@kth.se
[@thvasilo](#)

References

- ApproximateQueries on Flink:
<https://github.com/tlindener/ApproximateQueries/>
- Gelly Streaming
<https://github.com/vasia/gelly-streaming>
- Yahoo Dataskeches
<https://dataskeches.github.io/>
- Collection of links on streaming algorithms and sketches
<https://gist.github.com/debasishg/8172796>